

# A

DRAFT → CRITIQUE → REVISE

## Prompts write. Workflows hand off.

A reusable, source-bounded Claude Project workflow turns a fact ledger into concise portfolio copy, attacks its own claims, revises them, and stops at a human verification gate.

### 5 steps

Defined handoffs

### 5 runs

Real project inputs

### 27.7 sec

Observed AI average

### 1 gate

Human before publish

#### 01 · AUDIT ANCHOR

### One repeated writing task became a system.

My Week 1 audit identified “drafting narrative text or documentation for research papers” as a recurring task suited to collaboration with AI. The workflow preserves that collaboration: AI structures and challenges the draft; I remain accountable for facts, permissions, and publication.

**Tool** Claude Project with structured instructions

**Backend** None

**Trigger** RUN DCR WORKFLOW

**Auto-publish** Never

#### 02 · STEP DIAGRAM

### Each stage leaves an inspectable artifact.

1

#### Gather

Input → fact ledger



2

### Draft

Facts → cited copy



3

### Critique

Copy → issue log



4

### Revise

Issues → clean draft



5

### Verify

Draft → status + gate

**Closed fact boundary:** every checkable claim must trace to a numbered [F#] fact.  
Missing evidence becomes [NEEDS HUMAN FACT] ; it is never guessed.

## 03 · CONFIGURATION

# The whole build is portable, not hidden in a chat.

### Project instructions

Role, voice, five-step protocol, citation rule, output schema, and stop conditions.

### Open full instructions

### **Input contract**

Audience, purpose, numbered facts, constraints, and requested output format.

### **Open template**

### **Run controls**

Fresh-chat procedure, timer boundaries, and mandatory human checks.

### **Open checklist**

## **04 · FIVE REAL RUNS**

# **The critic caught a different risk each time.**

### **01 · DG CEMENT**

## **Offline-first field case**

Caught an implied sync algorithm that the source facts did not establish.

Human check: endpoint count and outage-loss evidence

**Full run** →

### **02 · AI CORE**

## **Text-to-SQL internship**

Prevented "secure" and "complete" from overstating lexical controls.

Human check: test totals and internship disclosure

**Full run** →

### **03 · AURA**

## **Agent-state debugging**

Refused to invent the exact root cause or unrecorded fix.

Human check: causal account and usage scope

**Full run** →

#### 04 · DATASAGE

### Multi-agent analysis

Scoped the speed and stability claims to the tested public data.

Human check: dataset names and benchmark method

**Full run** →

#### 05 · BLIND ACCEPTANCE

### Inventory system

Ran from a new input and removed an inferred user-interface flow.

Human check: whether alerts trigger automatic restocking

**Full run** →

#### 05 · HONEST TIME ACCOUNTING

## Measured generation; benchmark still open.

#### OBSERVED

**2m 18.7s**

Batch artifact generation for all five runs.

#### OBSERVED AVERAGE

**27.7s**

AI generation per run; excludes human review.

#### STILL TO MEASURE

## Manual baseline

Sara must time one equivalent draft and the live review pass.

Illustration only: if manual writing takes 20 minutes and review/correction takes 4 minutes after 28 seconds of generation, the workflow takes about 4.5 minutes and saves about 15.5 minutes per run before setup. This is a conditional estimate, not a measured personal claim. **See timestamps, formula, and fillable benchmark** →

06 · FAILURE MAP

Automation ends where responsibility begins.

| FAILURE                          | WORKFLOW RESPONSE          | HUMAN MUST CHECK                       |
|----------------------------------|----------------------------|----------------------------------------|
| Thin or ambiguous facts          | Marks [NEEDS HUMAN FACT]   | Supply evidence or cut the claim       |
| True fact, misleading emphasis   | Critique flags distortion  | Judge context and proportionality      |
| Numbers, names, dates            | Traces them to [F#]        | Verify against the original source     |
| Confidential internship material | Cannot infer permission    | Confirm disclosure and redact          |
| Plausible but invented detail    | Rejects unsupported claims | Inspect the final against evidence     |
| Tone that is technically valid   | Applies a voice card       | Decide whether it still sounds like me |

## Built, rerun live, and ready for human release.

✓ Three-plus distinct steps and explicit handoffs

---

✓ Five documented real inputs and outputs

---

✓ Brand-new input acceptance run

---

✓ Live Claude Project run completed

---

✓ Known failure points and human review named

---

△ Personal manual benchmark remains optional evidence

---

The workflow package is complete and reproducible. Sara confirms that she configured it in Claude and completed the live end-to-end run. The measured generation time remains separate from the conditional savings illustration because no personal manual baseline was supplied.